

# mdexplore

Fast Markdown explorer for Ubuntu/Linux desktop: browse `.md` files in a directory tree and preview fully rendered output instantly.

## User Guide

For a use-case/tutorial walkthrough, start with [USER-GUIDE.md](#).

## Companion App: pdfexplore

This repository also ships `pdfexplore`, a read-only PDF companion with the same shell philosophy as `mdexplore`.

- Primary docs: [pdfexplore/README.md](#)
- Maintainer guide: [pdfexplore/AGENTS.md](#)
- Launcher: `./pdfexplore.sh [PATH]`
- Its top-bar `Dark` button switches PDF pages to dark gray with light text and changes to `Light` while that mode is active.
- Its PDF preview mirrors mdexplore's navigation gutters with persistent-highlight markers on the left and active search-hit markers on the right, including qualifying-window behavior for `NEAR( . . . )` searches.
- Its extracted-text cache performs bounded garbage-collection passes on an independent 30-second cadence after 10 seconds of input idle, removing memory/disk entries and cache badges for PDFs that no longer exist independently of automatic prefetch. Filesystem validation and deletion stay on the idle worker and yield between entries when input resumes.
- It bounds live PDF WebEngine previews with a two-entry least-recently-used cache by default; evicted pages are discarded to prevent cumulative browser-memory growth.
- Automatic extracted-text prefetch resumes after 10 seconds idle, warms one PDF at a time with a one-second inter-batch cooldown, and moves cache probes off the GUI thread. Input across native controls and inside pdf.js cancels queued work immediately, and running extraction yields between PDF pages. Search extraction remains single-threaded and debounced.
- Its viewer bridge filters bridge-owned DOM mutations, coalesces marker publications, ignores identical overlay payloads, and avoids full-document overlay/index work during ordinary scrolling while preserving pdf.js's viewport-bounded scroll container and render virtualization.

## Running Multiple pdfexplore Instances

Multiple `pdfexplore` windows may run as separate OS processes, including on the same PDF tree.

- Color, view-session, and persistent-highlight sidecars use stable companion lock files and atomic replacement. A normal edit merges only the changed PDF entry with the latest sidecar on disk, so simultaneous edits to different PDFs are retained. Same-PDF highlight range edits are applied transactionally to the latest committed list; same-PDF color/view-session values use last-commit-wins semantics.
- Failed transactional highlight writes leave the disk payload unchanged, restore viewer/cache state from that payload, and report the save failure.
- Config writes replay this instance's newly recorded recent-root events over the latest on-disk history. The `default_root` is intentionally last-successful-writer-wins, while a failed non-blocking lock attempt leaves the existing config untouched for a later save attempt.
- The extracted-text disk cache uses short shared/exclusive transactions for opening entries, atomic commits, per-entry trimming, and garbage collection; gzip decompression and quota scans run outside the global lock. Path-stable, hash-striped producer locks also prevent two instances extracting the same source concurrently, including across source-file replacements. Each commit publishes a shared trim-dirty marker so any idle instance can enforce the quota without relying

on another process's in-memory counters. A shared activity heartbeat is written asynchronously and coalesced, and its timestamp is polled on a separate worker, so GUI input and idle schedulers never wait on cache-directory I/O; it delays idle work when any instance is active. A non-blocking background lease permits only one instance at a time to prefetch or collect.

- `pdfexplore.sh` serializes only virtual-environment/dependency bootstrap work with `flock` ; it releases that lock before Qt starts. PDF previews use Qt's off-the-record WebEngine profile, avoiding a shared persistent Chromium profile lock between processes.
- Sidecar changes are not pushed live between windows. Use `Refresh` or `F5` to invalidate disk-backed color/view/highlight snapshots, rescan markers, and reload the current PDF's persisted highlights while preserving tree expansion and selection where possible. A currently open window's live view tabs remain authoritative until their normal save/switch lifecycle.

Runtime settings are JSON-externalized:

- global/shared settings: `mdexplore.settings.json`
- pdfexplore-specific settings: `pdfexplore.settings.json`

The global file owns shared defaults used by both apps (colors, zoom defaults, print-layout constants, shared tokens), while the pdfexplore file owns only pdfexplore-specific behavior (worker/timer tuning, file names, viewer-bridge runtime settings).

## Settings Cheat Sheet

| What you want to change                                                   | File                                     | Jump to details                                             |
|---------------------------------------------------------------------------|------------------------------------------|-------------------------------------------------------------|
| Performance and responsiveness (worker counts, cache caps, timer cadence) | <a href="#">mdexplore.settings.json</a>  | <a href="#">Thread pools, cache caps, and timer cadence</a> |
| Rendering and preview behavior (zoom, highlight visuals)                  | <a href="#">mdexplore.settings.json</a>  | <a href="#">Preview highlights and zoom</a>                 |
| Search behavior and NEAR window tuning                                    | <a href="#">mdexplore.settings.json</a>  | <a href="#">Core paths and search tokens</a>                |
| Sidecar/session persistence keys and limits                               | <a href="#">mdexplore.settings.json</a>  | <a href="#">mdexplore app/session and sidecar keys</a>      |
| PDF print geometry and diagram fit heuristics                             | <a href="#">mdexplore.settings.json</a>  | <a href="#">PDF print geometry and fitting heuristics</a>   |
| Auto-calculated thread/cache behavior (advanced)                          | <a href="#">mdexplore.settings.json</a>  | <a href="#">Derived and calculated behaviors</a>            |
| Full mdexplore settings catalog                                           | <a href="#">mdexplore.settings.json</a>  | <a href="#">Externalized Settings Reference</a>             |
| pdfexplore-only runtime tuning                                            | <a href="#">pdfexplore.settings.json</a> | <a href="#">pdfexplore Runtime Settings</a>                 |

## Code Layout

The application still uses `mdexplore.py` as the main endpoint and primary UI module, but low-risk support code is now split into `mdexplore_app/` :

- `mdexplore.py` : main window, markdown renderer, preview/PDF orchestration, and app endpoint.
- `mdexplore_app/constants.py` : shared runtime/render constants.
- `mdexplore_app/runtime.py` : runtime environment helpers (config path, default root, GPU/software fallback, print layout knobs).
- `mdexplore_app/search.py` : extracted search tokenization, boolean/NEAR parsing, and per-file hit counting.
- `mdexplore_app/templates.py` : extracted HTML template-asset registry/renderer for preview document shells.
- `mdexplore_app/pdf.py` : PDF footer stamping, blank-page suppression, and PlantUML stderr formatting.
- `mdexplore_app/icons.py` : icon loading/recoloring helpers.
- `mdexplore_app/file_coordination.py` : stable advisory locks, atomic text-file replacement, and merge-safe per-file sidecar updates shared with `pdfexplore` .

- `mdexplore_app/tree.py` : markdown tree model and delegate ( `ColorizedMarkdownModel` , `MarkdownTreeItemDelegate` ).
- `mdexplore_app/tabs.py` : custom multi-view tab bar ( `ViewTabBar` ).
- `mdexplore_app/workers.py` : background worker classes for preview render, PlantUML, and PDF write/stamp steps.
- `mdexplore_app/fast_base64.py` : shared BASE64 helpers with vendor `fastbase64` + `pybase64` acceleration and `stdlib` fallback.

This is intentionally a first-stage modularization. Most behavior and call flow still lives in `mdexplore.py` so feature risk stays low while the file is gradually decomposed.

## Features

- Expandable left-pane directory tree rooted at a chosen folder.
- Shows Markdown files only ( `*.md` ), while still allowing folder navigation.
- Large right-pane rendered preview with scroll support.
- PlantUML diagram renders run in background workers so markdown preview stays responsive.
- Supports:
  - CommonMark + tables + strikethrough.
  - Markdown rendering defaults to `cmarkgfm` for speed, with automatic fallback to `markdown-it-py` when unavailable/incompatible.
  - TeX/LaTeX math via MathJax.
  - Mermaid diagrams with improved dark-theme contrast.
  - Optional Rust Mermaid backend via `mmdr ( --mermaid-backend rust )`.
  - PlantUML diagrams (asynchronous local render with placeholders).
  - Markdown callouts ( `> [!NOTE]` , `> [!TIP]` , `> [!IMPORTANT]` , `> [!WARNING]` , `> [!CAUTION]` ).
- Top actions:
  - **Recent** opens a dropdown of up to 35 retained root directories, with the entire list shown most-recent-first. The list refreshes from disk each time the menu is opened so multiple running instances stay in sync.  
A root is added only after it has been active for at least 30 seconds and then you navigate to another root.
  - **^** moves root up one directory level.
  - **Refresh** rescans the current directory view to pick up new/deleted files.
  - **PDF** exports the current preview to `<filename>.pdf` with centered page numbering ( `N of M` ).  
During export, retrievable image sources are inlined as BASE64 data URIs so images render in the PDF instead of relying on external links; broken links are left unchanged.
  - **Add View** creates another tabbed view of the same document at the current top visible line.
  - **Edit** opens the selected file in MarkText ( `/usr/bin/marktext` ).
- Window title shows the current effective root path.
- Effective-root directory row is always bold:
  - aqua-blue ( `#7fdfe8` ) when no active search matches are under that scope,
  - dark orange ( `#d79232` ) when active search has matches under that scope.
- Every directory containing matched descendant files uses dark-orange label text and shows an appended lighter-orange matching-file-count pill ( `1..99` , then `++` ), independent of selection. Pills update progressively during a scan and aggregate under the visible tree location of symlinked files/directories.
- Preview cache keyed by file timestamp and size for fast re-open.
- Inline preview BASE64 images are materialized in parallel per document (deduped by payload hash) to reduce first-load latency on image-heavy files.
- Copy-time BASE64 image conversion warms image targets in parallel before rewrite for faster large-file copy/export workflows.
- Mermaid SVGs are cached in-memory per app run to avoid re-rendering when returning to previously viewed files.
- Mermaid keeps separate in-memory cache modes for GUI ( `auto` ) and PDF ( `pdf` ) rendering.
- Navigating back to a cached file still performs a fresh stat check; changed files re-render automatically.
- Startup-generated icon assets are persisted under `~/.cache/mdexplore/icon-cache` :
  - app icon normalization output (from `mdexplor-icon.png` /fallback icon files),
  - two-tone icon recolor outputs used by the tab/tree UI.These are regenerated automatically when source asset identity (path/mtime/size) or render parameters change.
- **F5** refresh shortcut for directory view rescan (same behavior as **Refresh** button).
- **Ctrl++** / **Ctrl+-** / **Ctrl+0** zoom only the preview pane content in/out/reset.
- Preview-only zoom changes briefly show a percentage badge at the top of the preview pane.
- If the currently previewed markdown file changes on disk, preview auto-refreshes and shows a status bar message.
- Status bar reports active long-running work (preview load/render, PlantUML progress, PDF export) and returns to **Ready** instead of staying blank.
- Manual tree/preview pane resizing is preserved across **^** root navigation for the current app run.

- Right-click a Markdown file to assign a highlight color in the tree.
- Highlight colors persist per directory in `.mdexplore-colors.json` files.
- Available file-highlight colors are: Yellow, Green, Blue, Orange, Purple, Light Gray, Medium Gray, and Red.
- View-tab state persists per directory in `.mdexplore-views.json` files for documents that have more than one saved view or a custom tab label.
- Persistent preview text highlights are stored per directory in `.mdexplore-highlighting.json`.
- Files in the tree show a small light-gray tab badge when they currently have more than the default single view.
- Files with persistent preview text highlights show a separate purple marker badge in the tree.
- Markdown tree gutter badges are packed together in a fixed-width strip so the markdown file icon and filename stay aligned even when only some badges apply.
- Right-click tree menu includes **Clear in Directory** (non-recursive) and **Clear All** (recursive) to remove persisted file-highlight colors, and both actions prompt for confirmation before clearing.
- Top-right copy controls are labeled **Copy to: ( ) Clipboard ( ) Directory**, with **Clipboard** selected by default.
- A BASE64 image toggle button sits immediately to the left of **Copy to: :**
  - off tooltip: Turn BASE64 image encoding on
  - on tooltip: Turn BASE64 image encoding off
  - toggle state persists in `~/mdexplore.cfg` across restarts (`copy_base64_images_enabled`).
  - when enabled, copied markdown (clipboard staging or directory copy) converts retrievable image links (local paths and URLs) into embedded BASE64 `data: URIs`.
- In **Clipboard** mode, color buttons copy matching highlighted files and the pin button copies the currently previewed markdown file.
- In **Directory** mode, pin/color actions open a target-folder picker, then copy the file(s) into that folder.
- Directory copy also writes merged metadata entries for copied files into the target folder sidecars: `.mdexplore-colors.json`, `.mdexplore-views.json`, and `.mdexplore-highlighting.json`.
- Search box includes an explicit **X** clear control that clears the query and removes match styling.
- Active search evaluates the markdown files currently visible in the tree (root + expanded branches), and reruns automatically when tree visibility changes (expand/collapse/root/scope navigation).
- Search-hit files in the tree show a yellow hit-count pill in the left gutter.
- Matching filenames are styled bold+italic, and filename-term matches are rendered in yellow text.
- When search is active and a matched file is opened, preview matches are highlighted in yellow and the view scrolls to the first match.
- While dragging the preview scrollbar, mdexplore shows an approximate **current line / total lines** indicator beside the scrollbar handle.
- When search is active, preview scrollbar markers show the vertical positions of highlighted hits; clicking a marker jumps to the nearest hit in that cluster.
- Preview scroll position is remembered per markdown file for the current app session.
- Right-click selected text in the preview pane to use:
  - **Copy Rendered Text** for plain rendered selection text.
  - **Copy Source Markdown** for markdown source content.
 Copies matching source markdown using direct range mapping first, then selected-text/fuzzy line matching as fallback, and finally the full source file if no match is possible.

- Clipboard copy uses file URI MIME formats compatible with Nemo/Nautilus paste.
- Last effective root plus recent-root history are persisted to `~/.mdexplore.cfg` on root navigation and on exit.
  - Payload format is JSON with keys `default_root`, `recent_roots`, and `copy_base64_images_enabled` (legacy plain-text config is still accepted on read).
  - Recent roots are capped at 35 entries in rolling newest-first storage, and the entire menu retains that newest-first order.
  - Config writes use a lock file ( `~/.mdexplore.cfg.lock` ) with non-blocking, momentary locking.
  - If another instance holds the lock during a save attempt, that save is skipped silently.
  - Lock files older than 2 minutes are cleaned up automatically (silently).
  - If no directory is selected at quit time, the most recently selected/expanded directory is used for `default_root`.

## Externalized Settings Reference

`mdexplore` runtime constants are externalized in `mdexplore.settings.json`. `mdexplore_app/constants.py` is the loader and type gate for this file.

### Loading and override behavior

- The loader reads `mdexplore.settings.json` at startup.
- If the file is missing, unreadable, or not valid JSON, built-in defaults are used.
- Only known keys are consumed; unknown keys are ignored.
- Each key is type-checked against the built-in schema ( `str`, `int/float`, `list` ).
- Numeric keys accept integer/float JSON values and are cast to the target runtime type.

### Units and naming conventions

- `_MS` : milliseconds.
- `_SECONDS` : seconds.
- `_BYTES` : bytes.
- `_PX` : CSS/layout pixels.
- `_PT` : print font points.
- Ratio keys (for example `*_ASPECT_RATIO` ) are unitless floats.

## Key groups and defaults

### Core paths and search tokens

| Key                         | Default                                                                                                               | Purpose                                                 |
|-----------------------------|-----------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------|
| CONFIG_FILE_NAME            | .mdexplore.cfg                                                                                                        | User config filename under home directory.              |
| UI_ASSET_DIR_RELATIVE       | assets/ui                                                                                                             | Relative path for UI/icon assets.                       |
| JS_ASSET_DIR_RELATIVE       | assets/js                                                                                                             | Relative path for preview/PDF JS templates.             |
| TEMPLATE_ASSET_DIR_RELATIVE | assets/templates                                                                                                      | Relative path for HTML templates.                       |
| SEARCH_CLOSE_WORD_GAP       | 50                                                                                                                    | Word-distance threshold used by NEAR( . . . ) matching. |
| SEARCH_HIT_COUNT_FONT_PATHS | assets/ui/LiberationSansNarrow-Regular.ttf ,<br>/usr/share/fonts/truetype/liberation/LiberationSansNarrow-Regular.ttf | Ordered fallback list for tree hit-count pill font.     |

## Diagram state and restore controls

| Key                                 | Default                               | Purpose                                                         |
|-------------------------------------|---------------------------------------|-----------------------------------------------------------------|
| PDF_LANDSCAPE_PAGE_TOKEN            | MDEXPLORE_LANDSCAPE_PAGE_TOKEN        | Marker token used in print layout processing.                   |
| PDF_EXPORT_PRECHECK_MAX_ATTEMPTS    | 60                                    | Max precheck polling attempts before PDF export proceeds/fails. |
| PDF_EXPORT_PRECHECK_INTERVAL_MS     | 140                                   | Delay between export precheck probes.                           |
| MERMAID_CACHE_JSON_TOKEN            | __MDEXPLORE_MERMAID_CACHE_JSON__      | Placeholder token for Mermaid cache serialization.              |
| DIAGRAM_VIEW_STATE_JSON_TOKEN       | __MDEXPLORE_DIAGRAM_VIEW_STATE_JSON__ | Placeholder token for diagram view-state serialization.         |
| MERMAID_SVG_CACHE_MAX_ENTRIES       | 256                                   | Max in-memory Mermaid SVG cache entries.                        |
| MERMAID_SVG_MAX_CHARS               | 250000                                | Maximum Mermaid source size accepted for caching.               |
| PLANTUML_RESTORE_BATCH_SIZE         | 2                                     | Number of PlantUML restore operations processed per batch.      |
| MERMAID_CACHE_RESTORE_BATCH_SIZE    | 2                                     | Number of Mermaid restore operations processed per batch.       |
| RESTORE_OVERLAY_TIMEOUT_SECONDS     | 25.0                                  | Upper bound for restore overlay lifecycle.                      |
| RESTORE_OVERLAY_SHOW_DELAY_MS       | 350                                   | Delay before showing restore overlay.                           |
| RESTORE_OVERLAY_MAX_VISIBLE_SECONDS | 1.0                                   | Cap on visible overlay duration.                                |
| MERMAID_BACKEND_JS                  | js                                    | Canonical value for JS Mermaid backend selector.                |
| MERMAID_BACKEND_RUST                | rust                                  | Canonical value for Rust Mermaid backend selector.              |



## Preview highlights and zoom

| Key                                                 | Default                   | Purpose                                                 |
|-----------------------------------------------------|---------------------------|---------------------------------------------------------|
| PREVIEW_PERSISTENT_HIGHLIGHT_COLOR                  | rgba(102, 86, 178, 0.36)  | Default persistent preview highlight color.             |
| PREVIEW_PERSISTENT_HIGHLIGHT_IMPORTANT_COLOR        | rgba(225, 214, 255, 0.76) | Default persistent important-highlight color.           |
| PREVIEW_PERSISTENT_HIGHLIGHT_IMPORTANT_TEXT_COLOR   | #170534                   | Foreground text color used on important highlights.     |
| PREVIEW_PERSISTENT_HIGHLIGHT_MARKER_COLOR           | rgba(112, 90, 188, 0.92)  | Scroll-marker color for normal highlights.              |
| PREVIEW_PERSISTENT_HIGHLIGHT_IMPORTANT_MARKER_COLOR | rgba(154, 132, 220, 0.96) | Scroll-marker color for important highlights.           |
| PREVIEW_HIGHLIGHT_KIND_NORMAL                       | normal                    | Internal key name for normal highlight kind.            |
| PREVIEW_HIGHLIGHT_KIND_IMPORTANT                    | important                 | Internal key name for important highlight kind.         |
| PREVIEW_ZOOM_STEP                                   | 0.1                       | Zoom increment/decrement step per action.               |
| PREVIEW_ZOOM_MIN                                    | 0.35                      | Minimum preview zoom factor.                            |
| PREVIEW_ZOOM_MAX                                    | 3.0                       | Maximum preview zoom factor.                            |
| PREVIEW_ZOOM_RESET                                  | 1.0                       | Zoom factor used for reset action.                      |
| PREVIEW_ZOOM_OVERLAY_TIMEOUT_MS                     | 850                       | Duration of transient zoom-percent overlay.             |
| PREVIEW_SETHTML_MAX_BYTES                           | 650000                    | Guard rail for <code>setHtml</code> payload size paths. |

## SVG icon rendering controls

| Key                        | Default | Purpose                                              |
|----------------------------|---------|------------------------------------------------------|
| SVG_ICON_RENDER_OVERSAMPLE | 4       | Oversample multiplier for SVG rasterization quality. |
| SVG_ICON_ALPHA_CUTOFF      | 1       | Alpha threshold used during icon post-processing.    |

## PDF print geometry and fitting heuristics

| Key                                       | Default | Purpose                                                   |
|-------------------------------------------|---------|-----------------------------------------------------------|
| MIN_PRINT_DIAGRAM_FONT_PT                 | 3       | Minimum diagram font size used by print fitter.           |
| MAX_PRINT_DIAGRAM_FONT_PT                 | 12.0    | Maximum diagram font size used by print fitter.           |
| PDF_PRINT_HEADING_TO_DIAGRAM_GAP_PX       | 16      | Vertical spacing between heading block and diagram.       |
| PDF_PRINT_LAYOUT_SAFETY_PX                | 40      | Extra padding to avoid clipping at boundaries.            |
| PDF_PRINT_KEEP_MIN_HEIGHT_PX              | 120     | Minimum element height considered keepable without split. |
| PDF_PRINT_PORTRAIT_LETTER_WIDTH_PX        | 1000    | Portrait page model width for fitting math.               |
| PDF_PRINT_PORTRAIT_LETTER_HEIGHT_PX       | 1310    | Portrait page model height for fitting math.              |
| PDF_PRINT_LANDSCAPE_LETTER_WIDTH_PX       | 1100    | Landscape page model width for fitting math.              |
| PDF_PRINT_LANDSCAPE_LETTER_HEIGHT_PX      | 860     | Landscape page model height for fitting math.             |
| PDF_PRINT_HORIZONTAL_MARGIN_PX            | 80      | Horizontal margin budget in print layout solver.          |
| PDF_PRINT_VERTICAL_MARGIN_PX              | 130     | Vertical margin budget in print layout solver.            |
| PDF_PRINT_PORTRAIT_MIN_WIDTH_PX           | 320     | Minimum accepted portrait diagram width.                  |
| PDF_PRINT_PORTRAIT_MIN_HEIGHT_PX          | 320     | Minimum accepted portrait diagram height.                 |
| PDF_PRINT_LANDSCAPE_MIN_WIDTH_PX          | 360     | Minimum accepted landscape diagram width.                 |
| PDF_PRINT_LANDSCAPE_MIN_HEIGHT_PX         | 260     | Minimum accepted landscape diagram height.                |
| PDF_PRINT_WIDE_DIAGRAM_ASPECT_RATIO       | 1.19    | Aspect-ratio threshold used to classify diagrams as wide. |
| PDF_PRINT_WIDE_DIAGRAM_LANDSCAPE_GAIN     | 1.06    | Gain factor favoring landscape for wide diagrams.         |
| PDF_PRINT_PLANTUML_LANDSCAPE_ASPECT_RATIO | 1.18    | PlantUML-specific landscape threshold.                    |

## mdexplore app/session and sidecar keys

| Key                                             | Default                                                           | Purpose                                                                 |
|-------------------------------------------------|-------------------------------------------------------------------|-------------------------------------------------------------------------|
| MDEXPLORE_MAX_DOCUMENT_VIEWS                    | 8                                                                 | Maximum tabs/views per document.                                        |
| MDEXPLORE_MAX_RECENT_ROOT_DIRECTORIES           | 35                                                                | Persistent recent-root history capacity.                                |
| MDEXPLORE_RECENT_ROOT_MENU_MRU_COUNT            | 10                                                                | Legacy compatibility setting; the full Recent menu is now newest-first. |
| MDEXPLORE_MIN_RECENT_ROOT_DWELL_SECONDS         | 30.0                                                              | Minimum time a root must stay active before being persisted as recent.  |
| MDEXPLORE_CONFIG_LOCK_STALE_SECONDS             | 120.0                                                             | Age threshold for lock-file stale cleanup.                              |
| MDEXPLORE_VIEWS_FILE_NAME                       | .mdexplore-views.json                                             | Sidecar file for per-document multi-view sessions.                      |
| MDEXPLORE_HIGHLIGHTING_FILE_NAME                | .mdexplore-highlighting.json                                      | Sidecar file for persistent text highlights.                            |
| MDEXPLORE_CONFIG_DEFAULT_ROOT_KEY               | default_root                                                      | Config JSON key for default root path.                                  |
| MDEXPLORE_CONFIG_RECENT_ROOTS_KEY               | recent_roots                                                      | Config JSON key for recent-root list.                                   |
| MDEXPLORE_CONFIG_COPY_BASE64_IMAGES_ENABLED_KEY | copy_base64_images_enabled                                        | Config JSON key for BASE64 copy toggle state.                           |
| MDEXPLORE_DEBUG_LOG_FILE_NAME                   | mdexplore.log                                                     | Debug log filename under cache/log area.                                |
| MDEXPLORE_DEBUG_LOG_MAX_LINES                   | 10000                                                             | In-memory cap for retained debug log lines.                             |
| MDEXPLORE_HIGHLIGHT_COLORS                      | Yellow, Green, Blue, Orange, Purple, Light Gray, Medium Gray, Red | Default file-highlight color palette and order.                         |

## Thread pools, cache caps, and timer cadence

| Key                                                  | Default    | Purpose                                                                                                     |
|------------------------------------------------------|------------|-------------------------------------------------------------------------------------------------------------|
| MEXPLORE_DEFAULT_SEARCH_SCAN_MAX_THREADS             | 0          | Search worker thread override. 0 means auto (see formula below).                                            |
| MEXPLORE_SEARCH_WORKER_CHUNK_SIZE                    | 8          | Markdown files per search worker batch; small batches preserve prompt progressive results and cancellation. |
| MEXPLORE_SEARCH_WORKER_YIELD_SECONDS                 | 0.001      | Cancellation-aware pause between searched files that releases the GIL for the GUI.                          |
| MEXPLORE_DEFAULT_BASE64_IMAGE_WORKER_THREADS         | 0          | BASE64 materialization worker override. 0 means auto (see formula below).                                   |
| MEXPLORE_PREVIEW_HTTP_CACHE_MAX_BYTES                | 4294967296 | Desired per-instance preview HTTP cache cap before clamping.                                                |
| MEXPLORE_PREVIEW_HTTP_CACHE_MAX_BYTES_QT_MAX         | 2147483647 | Qt-safe upper cap applied to HTTP cache size.                                                               |
| MEXPLORE_PREVIEW_HTTP_CACHE_MIN_BYTES                | 67108864   | Lower bound for reduced cache cap.                                                                          |
| MEXPLORE_PREVIEW_HTTP_CACHE_FREE_SPACE_RESERVE_BYTES | 536870912  | Free-space reserve that must remain after cache assignment.                                                 |
| MEXPLORE_PREVIEW_WEBENGINE_INSTANCE_STALE_SECONDS    | 172800     | Stale age for preview profile cleanup (48h).                                                                |
| MEXPLORE_INLINE_DATA_IMAGE_POOL_THREADS              | 1          | Worker pool size for inline data-image tasks.                                                               |
| MEXPLORE_TREE_MARKER_SCAN_POOL_THREADS               | 1          | Worker pool size for tree marker scans.                                                                     |
| MEXPLORE_PLANTUML_POOL_MAX_THREADS                   | 6          | PlantUML render worker max threads.                                                                         |
| MEXPLORE_PDF_POOL_MAX_THREADS                        | 1          | PDF write/stamp worker max threads.                                                                         |
| MEXPLORE_MATCH_TIMER_INTERVAL_MS                     | 3000       | Search trigger timer cadence.                                                                               |
| MEXPLORE_TREE_SEARCH_REFRESH_DEBOUNCE_MS             | 180        | Debounce delay for tree-change search reruns.                                                               |
| MEXPLORE_SEARCH_PROGRESS_PUBLISH_INTERVAL_MS         | 100        | Maximum coalescing interval for progressive search-pill updates after the first immediate hit.              |
| MEXPLORE_SCROLL_CAPTURE_INTERVAL_MS                  | 200        | Scroll state capture timer cadence.                                                                         |
| MEXPLORE_DIAGRAM_STATE_CAPTURE_INTERVAL_MS           | 250        | Diagram state snapshot timer cadence.                                                                       |
| MEXPLORE_PREVIEW_ACTION_POLL_INTERVAL_MS             | 220        | Preview action/pending-state poll cadence.                                                                  |

| Key                                                     | Default           | Purpose                                         |
|---------------------------------------------------------|-------------------|-------------------------------------------------|
| <code>MDEXPLORE_STATUS_IDLE_INTERVAL_MS</code>          | <code>900</code>  | Idle-status update cadence.                     |
| <code>MDEXPLORE_FILE_CHANGE_WATCH_INTERVAL_MS</code>    | <code>1200</code> | Poll interval for source-file change detection. |
| <code>MDEXPLORE_RESTORE_OVERLAY_POLL_INTERVAL_MS</code> | <code>170</code>  | Overlay lifecycle poll cadence.                 |
| <code>MDEXPLORE_PDF_DIAGRAM_PROBE_INTERVAL_MS</code>    | <code>220</code>  | PDF diagram readiness probe cadence.            |

## Derived and calculated behaviors

- `MDEXPLORE_DEFAULT_SEARCH_SCAN_MAX_THREADS`
  - if `> 0` : used as-is.
  - if `0` : uses one worker. Markdown matching is regex/GIL-heavy, so additional Python threads can make searches slower while starving the GUI.
- `MDEXPLORE_DEFAULT_BASE64_IMAGE_WORKER_THREADS`
  - if `> 0` : used as-is.
  - if `0` : computed as `max(2, min(24, (cpu_count or 2) * 2))`.
- Preview HTTP cache cap
  - initial target is `min(MDEXPLORE_PREVIEW_HTTP_CACHE_MAX_BYTES, MDEXPLORE_PREVIEW_HTTP_CACHE_MAX_BYTES_QT_MAX)`.
  - with multiple live mdexplore processes, cache cap can be halved repeatedly until free-space reserve and minimum bounds are both satisfied.

## Practical tuning notes

- Prefer changing one setting at a time and validating with one known-problem and one known-good document.
- Keep related values consistent (for example zoom min/max/reset and search debounce/timer pairs).
- For portability, keep absolute paths ( `SEARCH_HIT_COUNT_FONT_PATHS` ) as fallbacks after repository-relative entries.
- When changing sidecar/config key names, migrate existing files or preserve backward compatibility.

## Requirements

- Ubuntu/Linux desktop with GUI.
- Python 3.10+.
- `python3-venv` package available.
- Internet access for Mermaid only when no local Mermaid bundle is available.
- Internet access for MathJax only when no local MathJax bundle is available.
- Java runtime ( `java` in `PATH` ) for local PlantUML rendering.
- `plantuml.jar` available (vendored path by default, or set `PLANTUML_JAR` ).
- Optional: MarkText installed at `/usr/bin/marktext` for Edit.
- Optional: `mmdr` in `PATH` (or `MDEXPLORE_MERMAID_RS_BIN` ) for Rust Mermaid backend.

## Quick Start

For a new Debian/Ubuntu host, run the host bootstrap once:

```
/path/to/mdexplore/setup-host.sh
```

`setup-host.sh` installs/verifies the native host requirements (including Python venv support, Java, Poppler, and Tesseract) and then runs the project bootstrap. It is safe to rerun. Use `setup-host.sh --check-only` to verify an existing machine without making changes.

If the host dependencies are already installed, you can run only the project bootstrap:

```
/path/to/mdexplore/setup-mdexplore.sh
/path/to/mdexplore/mdexplore.sh
```

`setup-mdexplore.sh` is the project bootstrap path. It:

- creates or updates `.venv`
- installs Python dependencies from `requirements.txt`
- verifies tracked UI assets under `assets/ui`
- downloads vendored MathJax, Mermaid, and PlantUML runtime assets if they are missing
- downloads/builds the Rust Mermaid renderer under `vendor/mermaid-rs-renderer`

It is safe to rerun. After bootstrap, use `mdexplore.sh` for normal launches.

When no `PATH` is supplied, the app opens:

1. `default_root` from `~/mdexplore.cfg` (if valid), otherwise
2. your home directory.

To open a specific root directory:

```
/path/to/mdexplore/mdexplore.sh /path/to/notes
```

What the launcher does:

- Creates `.venv` inside the project if missing.
- Uses `.venv/bin/python` directly (does not alter your current shell session).
- Installs dependencies when `requirements.txt` changes.
- Stores a runtime import-check stamp in `.venv/.runtime-import-check.sha256`.
- Runs runtime import verification when one of these is true:
  - `.venv` is newly created,
  - `requirements.txt` hash changed,
  - runtime import-check stamp is missing or does not match requirements hash.
- If runtime import verification fails, self-heals by reinstalling dependencies.
- Runs the app.

## Usage

### Wrapper script

```
mdexplore.sh [--mermaid-backend js|rust] [PATH]
```

- `PATH` is optional.
- `--mermaid-backend` is optional. For `mdexplore.sh`, the launcher defaults to `rust` when not specified. ( `mdexplore.py` direct runs default to `js`.)
- Supports plain paths and `file://` URIs (for `.desktop %u` launches).
- If a file path is passed, `mdexplore` opens its parent directory.
- If omitted, `default_root` in `~/mdexplore.cfg` is used (falling back to home directory).

- `--help` prints usage.

## Direct Python run

```
python3 -m pip install -r /path/to/mdexplore/requirements.txt
python3 /path/to/mdexplore/mdexplore.py [--mermaid-backend js|rust] [PATH]
```

If `PATH` is omitted for direct run, the same config/home default rule applies.

## hfind CLI

`hfind` reuses mdexplore search syntax (AND/OR/NOT, quoted terms, `NEAR( . . . )` ) for file discovery.

Its hfind-specific runtime defaults are stored in `hfind.settings.json`, following the same JSON-backed settings pattern used by mdexplore and pdfexplore. The file includes the CPU target plus hfind-only worker, CPU sampling, binary detection, PDF/OCR settings, and Office extraction time/ZIP-safety limits. `HFIND_CPU_LIMIT` and `HFIND_SEARCH_THREADS` continue to override their corresponding defaults when set in the environment.

Run via wrapper:

```
./hfind.sh [--query QUERY|-q QUERY] [--base|-b] [--content|-c] [--recursive|-r] [--verbose|-v] [--pdf|-p] [--ocr-pdf|-o] [--wip|-w] [--exclude PATH|-e PATH] [--type TYPE|-t TYPE] [--links|-l] [--cpu-limit PERCENT] [--sort|-s] [--sort-case-sensitive|-S] PATTERN [PATTERN ...]
```

Run via Python directly:

```
python3 /path/to/mdexplore/hfind.py [--query QUERY|-q QUERY] [--base|-b] [--content|-c] [--recursive|-r] [--verbose|-v] [--pdf|-p] [--ocr-pdf|-o] [--wip|-w] [--exclude PATH|-e PATH] [--type TYPE|-t TYPE] [--links|-l] [--cpu-limit PERCENT] [--sort|-s] [--sort-case-sensitive|-S] PATTERN [PATTERN ...]
```

Notes:

- `--content / -c` is optional.
- `--recursive / -r` is optional.
- Default filename/query matching target is full discovered path (directories + filename).
- `--base / -b` switches query matching target to basename only (filename + extension).
- `--verbose / -v` prints matching line(s) under each matched file and highlights hit text in yellow.
- `--pdf / -p` enables searching extracted text inside `.pdf` files.
- `--ocr-pdf / -o` enables OCR fallback for PDFs that appear to be scans. It implies `--pdf`, but content search still requires `--content / -c`. The short form can be used alone ( `-o` ) or stacked with other short switches, such as `-cro` or `-crvo`.
  - Normal searchable PDFs stay on the fast text-extraction path.
  - PDFs with no extractable text, sparse text plus raster-image evidence, or mostly scan-like image pages are treated as OCR candidates.
  - When a mixed PDF already contains searchable text, that native text is preserved and combined with OCR output from its image pages for matching; OCR does not replace the text that was already extracted.
  - OCR uses Poppler `pdftoppm` plus Tesseract when both commands are available; if either is unavailable, hfind falls back to ordinary PDF text extraction.
- `--wip / -w` prints one gray progress line for every file as it finishes being checked. The progress line uses the file's full path and appends operations that were actually performed, such as `[content read]`, `[PDF text]`, `[OCR]`, or `[binary skipped]`.

- WIP lines are progressive even when `--sort / -s` or `--sort-case-sensitive / -S` is active.
  - Without sorting, matches still appear progressively and are then replayed together at the end so they are not lost among non-matching WIP lines. The replay uses the identical filepath colors, verbose excerpts, hit highlighting, and other formatting as the original match output.
  - With `-s / -S`, matches are retained and emitted together at the end in the requested order, using that same complete formatting.
  - A WIP line for a file with no match is never added to the final match output; only actual matches are replayed or sorted.
- `--exclude PATH / -e PATH` excludes that path and every file beneath it. The option is repeatable, and `--exclude=PATH / -e=PATH` are also accepted.
  - Exactly one path belongs to each `-e / --exclude` occurrence; the path is consumed atomically by the option and is never interpreted as the query or a search pattern/base path.
  - `~` is expanded using the current user's home directory, and relative exclude paths are resolved from the current working directory.
- `--type TYPE / -t TYPE` limits the search to files whose extension is explicitly named. Types are written without the initial period, matching is case-insensitive, and the option is repeatable: `-t pdf -t txt`.
  - Multiple types can also be supplied with a pipe, such as `-t "pdf|txt"`. Quote or escape the pipe so the shell passes it to `hfind` instead of treating it as a pipeline.
  - Each `-t / --type` consumes exactly one atomic TYPE value and does not treat that value as the query or a search pattern.
  - Images are deliberately opt-in. Specifying `png`, `jpg`, `jpeg`, `tif`, or `tiff` causes matching images to be OCR'd automatically during content searches; `-o` is not needed because it controls PDF OCR only. Without an image type in `-t / --type`, images are skipped.
  - The pseudo-type `all-images` expands to every supported image type (`png`, `jpg`, `jpeg`, `tif`, and `tiff`).
  - Microsoft Word, Excel, and PowerPoint files are content-searchable by default and can also be selected explicitly. Supported modern formats are `docx`, `docm`, `dotx`, `dotm`, `xlsx`, `xlsm`, `xltx`, `xltm`, `pptx`, `pptm`, `ppsx`, `ppsm`, `potx`, and `potm`. Their document XML is read directly from the ZIP package with member-count and expanded-size safety limits.
  - Supported legacy formats are `doc`, `dot`, `xls`, `xlt`, `ppt`, and `pps`. They use `antiword / catdoc`, `xls2csv`, or `catppt`; `setup-host.sh` installs and verifies those tools.
  - The pseudo-type `all-office` expands to every supported legacy and modern Microsoft Office type.
  - Pseudo-types can be repeated or combined with pipes like ordinary types, including `-t=all-images|all-office` (quote or escape the pipe when required by your shell).
  - When no type option is supplied, `hfind` considers every matched non-image type; Office extraction is automatic when `-content / -c` is active.
- Symlinks are ignored by default. This includes symlink files and files that would only be reached by traversing a symlinked directory.
- `--links / -l` opts in to symlinks, allowing symlink files and recursive traversal through symlinked directories. The short flag may be stacked with other short flags, for example `-rl` or `-crvl`.
- `--cpu-limit PERCENT` dynamically throttles work based on whole-system CPU use. The default is `90`; `--cpu-limit 0` disables CPU throttling.
  - `HFIND_CPU_LIMIT` sets the default CPU target.
  - `HFIND_SEARCH_THREADS` remains the maximum worker-pool size; `hfind` may keep fewer workers active while the system is busy.
- Default output is progressive: matches print as they are found.
- Ctrl-C interruption exits with status `130` after the single `Search interrupted by user.` message. For this interrupt path, `hfind` bypasses Python's normal interpreter-shutdown callbacks so worker/thread/tempfile cleanup tracebacks are not dumped after the interruption.
- `--sort / -s` waits for full scan, then prints results sorted case-insensitively.
- `--sort-case-sensitive / -S` waits for full scan, then prints results sorted case-sensitively.
  - when either sort mode is enabled, `hfind` prints:  
One moment please... finding all matches to sort them



- `NEAR(...)` is strict in both `mdexplore` and `hfind`: a hit is only produced when all NEAR terms occur within the configured 50-word window.
- In `hfind`, `NEAR(...)` uses the active stream: default path text, basename text with `-b`, and path/base plus file content with `-c`.
- For content search, inline `data:image/...;base64,...` payloads are ignored in both `mdexplore` and `hfind` to avoid false positives from embedded image data.
- In verbose `NEAR(...)` output, numbering reflects match semantics:
  - lines that independently satisfy `NEAR(...)` are each numbered,
  - lines shown only as contiguous context for a cross-line NEAR window are grouped and only the first line is numbered.
- Quoted terms with exactly one leading and/or trailing literal space use boundary-aware matching at that edge (word boundary, punctuation, or line break can satisfy it).
  - Example: `'Nico '` can match `Nico.`, `Nico)`, `Nico;`, or `Nico` followed by a line break.
  - Example: `' Nico '` can match `\nNico\n` and `His name is Nico.`
  - Exception: this boundary rule applies only to a single leading/trailing space. Terms like `' Nico'` or `'Nico '` require those two literal spaces.
- Short flags are stackable in any order (for example `-cr`, `-rc`).
- If `-q / --query` is omitted, the first positional string is treated as the query.
- Base mode checks basename only (including extension, no path).

Examples:

```
./hfind.sh --query "OR(fred, paul)" --content --recursive *.txt
```

Recursively searches from current directory for readable `.txt` files whose path or content contains (case-insensitive) `fred` or `paul`.

```
./hfind.sh -q "OR(fred, paul)" -cr *.txt
./hfind.sh -cr "OR(fred, paul)" *.txt
./hfind.sh --recursive -c "OR(fred, paul)" *.txt
./hfind.sh -crvp "OR(fred, paul)" *.txt
./hfind.sh -crvps "OR(fred, paul)" *.txt
./hfind.sh -crvpS "OR(fred, paul)" *.txt
./hfind.sh -crvo "invoice" ".*/**/*.pdf"
./hfind.sh -crvw "invoice" ".*/**/*.pdf"
./hfind.sh -r -e ~/my-dir -e ~/my-dir-2 "invoice" "*.pdf"
./hfind.sh -rl "invoice" "*.pdf"
./hfind.sh -r -t pdf -t txt "invoice" ""
./hfind.sh -r -t "pdf|txt" "invoice" ""
./hfind.sh -cr -t "png|jpg|jpeg|tif|tiff" "invoice number" ""
./hfind.sh -cr -t "doc|docx|xls|xlsx|ppt|pptx" "roadmap" ""
./hfind.sh -cr -t=all-images "invoice number" ""
./hfind.sh -cr -t=all-office "roadmap" ""
./hfind.sh -cr -t="all-images|all-office" "confidential" ""
./hfind.sh -crv -o --cpu-limit 70 "invoice" ".*/**/*.pdf"
```

All above are valid shorthand forms.

```
./hfind.sh -rs "conflicted" ".*/**/*"
./hfind.sh -rS "conflicted" ".*/**/*"
```

Shows the two sort modes over recursive path matches:

- `-rs` sorts case-insensitively.
- `-rS` sorts case-sensitively.

```
./hfind.sh --query "OR(fred, paul)" *.txt
```

Searches discovered paths only (current directory, non-recursive).

```
./hfind.sh -rb "site" "**/*"
```

Recursively searches using basename-only matching (legacy filename behavior).

```
./hfind.sh -q "AND('Fred',paul)" /path/to/directory/*.md
```

Searches discovered paths in `/path/to/directory/` (non-recursive) where query requires case-sensitive `Fred` and case-insensitive `paul`.

```
./hfind.sh -rc "NEAR(\"Fred is my friend\", \"is a nice guy\")" "/path/to my/directory/*.md"
```

Uses escaped double quotes inside a `NEAR(...)` query while recursively searching matching markdown files.

```
./hfind.sh -crv "NEAR(alpha,beta)" "**/*.md"
```

For each matching file, prints matching line numbers and highlights the triggering terms in yellow. For `NEAR(...)`, this may include multiple nearby lines so both terms are visible.

```
./hfind.sh -cv "'Nico '" "**/*.md"
./hfind.sh -cv "' Nico '" "**/*.md"
```

Demonstrates the single-space boundary behavior for quoted terms (including punctuation and line-break boundaries).

## Full Bootstrap Script

```
setup-mdexplore.sh [--skip-python] [--skip-assets] [--skip-rust] [--rebuild-rust]
```

- `--skip-python` leaves the existing `.venv` untouched.
- `--skip-assets` skips vendored MathJax / Mermaid JS / PlantUML runtime asset checks.
- `--skip-rust` skips Rust Mermaid bootstrap/build.
- `--rebuild-rust` forces a fresh `cargo build --release --locked` for `mmdr`.

By default the script prefers the vendored `vendor/mermaid-rs-renderer/` checkout if it already exists. If it does not exist, the script clones it from:

```
https://github.com/ljehuang/mermaid-rs-renderer.git
```

If `cargo` is missing, the setup script installs a minimal Rust toolchain through `rustup` so it can build `mmdr` locally.

## Headless Regression Tests

Run the GUI regression suite under `xvfb` so `QWebEngineView` has a display:

```
xvfb-run -a .venv/bin/python -m unittest discover -s tests -v
```

The current suites in `tests/test_preview_regressions.py`,  
`tests/test_search_query_syntax.py`,  
`tests/test_pdf_layout_hints.py`,  
`tests/test_template_assets.py`,  
`tests/test_tab_bar_layout.py`,  
and `tests/test_window_layout.py`.

cover:

- saved view-tab round trips on the `test/testdoc.md` fixture,
- left-gutter named-view markers restoring the same saved positions as tabs,
- left-gutter persistent highlight markers navigating to visible highlights,
- right-gutter search markers navigating to visible search hits,
- same-document live preview search coverage for unquoted terms, double/single quotes, trailing-space phrases, implicit AND, function-style AND/OR, `NOT`, and `NEAR(...)` variants,
- search parser/matcher handling for unquoted terms, double-quoted phrases, single-quoted case-sensitive phrases, and literal apostrophes inside double-quoted phrases,
- PDF post-pass TOC detection and landscape/diagram page-flag classification,
- HTML preview-template asset loading and `MarkdownRenderer.render_document()` integration,
- custom tab-bar label budgeting and stale close-button geometry fallback.
- top-bar path-label layout protection so long document paths do not force the main window to stay overly wide.

## File Highlights

- In the file tree, right-click any `.md` file.
- Choose a highlight color ( `Highlight Yellow` , then `... <Color>` ), or clear it.
- Colors are persisted in `.mdexplore-colors.json` in each affected directory.
- Available colors include `Light Gray` and `Medium Gray` in addition to the original color set.
- If a directory is not writable, color persistence fails quietly by design.
- Use top-right `Copy to: () Clipboard () Directory` controls to choose copy destination mode.
- In `Clipboard` mode, color buttons copy files with a given highlight color and the pin button copies the currently previewed markdown file path payload.
- In `Directory` mode, pin/color actions open a folder chooser that defaults to: previously selected destination folder, else current effective root.
- Directory copy writes file content and merges copied-file metadata into target `.mdexplore-colors.json` , `.mdexplore-views.json` , and `.mdexplore-highlighting.json` (creating sidecars when missing, updating existing sidecars by filename when present).
- Right-click a directory (or file row) and use `Clear in Directory` for non-recursive clear in that directory, with confirmation.
- Right-click a directory (or file row) and use `Clear All` for recursive clear under that scope, with confirmation.
- Color-copy match collection is recursive and uses scope in this order: selected directory, else most recently selected/expanded directory, else root.

## Preview Text Highlights

- In the preview pane, right-click selected text and choose either `Highlight` or `Highlight Important` to add a persistent text highlight behind that rendered text.
- Normal highlights use a darker subdued purple. Important highlights use a lighter purple so the two categories are visually distinct on the dark theme.
- Selecting part or all of an existing highlighted block and applying the other action converts just that selected range to the chosen highlight type.
- Highlighted preview text also appears as left-side navigation markers in the preview gutter. Important-highlight markers are lighter than normal-highlight markers, and longer markers indicate highlights that span more lines.
- Clicking a highlight marker jumps to the corresponding highlighted block.
- Named views with a saved `Return to beginning` anchor also appear as color-matched left-gutter markers in the preview, layered above highlight markers when the two coincide.
- Right-click inside an existing highlighted block to remove it with `Remove Highlight` .
- If the current selection overlaps existing highlighted text and also includes unhighlighted text, both highlight actions and `Remove Highlight` remain available so the block can be extended, converted, or removed.
- Preview text highlights persist per directory in `.mdexplore-highlighting.json` .
- The tree also mirrors persisted preview highlights with a marker badge so highlighted documents remain easy to spot while browsing directories.
- The tree gutter badge order is:  
hit-count pill, highlight marker, views badge, markdown file icon.
- Highlight persistence is best-effort: if the directory is not writable, mdexplore fails quietly rather than interrupting preview use.

## Preview Zoom

- `Ctrl++` zooms the preview pane in.
- `Ctrl+-` zooms the preview pane out.
- `Ctrl+0` resets preview zoom to `100%`.
- These shortcuts affect only the preview content, not the tree pane, toolbar, or overall window layout.
- Each zoom change briefly shows a percentage badge at the top-center of the preview pane and also reports the same value in the status bar.

## PDF Export

- Click `PDF` (between `Refresh` and `Add View`) to export the currently previewed file.
- Output path is the previewed file path with `.pdf` extension in the same directory.
- Export is based on the active document content and print-prepared preview state (markdown + math + Mermaid + PlantUML).
- Export waits briefly for math/diagram/font readiness to reduce rendering artifacts.
- Mermaid PDF behavior is backend-specific:
  - JS backend: Mermaid is rendered through a print-safe monochrome/grayscale path.
  - Rust backend: Mermaid starts from a dedicated PDF SVG set rendered by `mmdx` (default theming), then applies print-safe grayscale normalization (multi-shade, with readable dark text).
- Export auto-scales page content into a print-style layout with top/side margins and an uncluttered footer band.
- Landscape diagram sections are isolated onto dedicated print blocks and use a tighter horizontal margin budget so wide diagrams can make use of the rotated page.
- Headed landscape sections are anchored to the page that actually contains the diagram heading/content, avoiding earlier prose pages being rotated instead.
- Pages that look like a table of contents are never promoted to landscape just because they mention a landscape section heading in the TOC text.
- Footer number font size is matched to the document's dominant scaled body text size.
- Pages are stamped with centered footer numbering as `1 of N`, `2 of N`, etc.

## PDF Layout Tuning

PDF pagination and diagram sizing are controlled by a small group of constants near the top of `mdexplore.py`. These are the first settings to adjust if PDF output starts making poor keep/spill/landscape choices.

Most important knobs:

- `MIN_PRINT_DIAGRAM_FONT_PT`
  - Lower bound for the largest font in a diagram when `mdexplore` tries to keep that diagram on a single page.
  - Lower it if you want tall diagrams to stay on one page more often.
  - Raise it if you prefer multi-page spill over small text.
- `MAX_PRINT_DIAGRAM_FONT_PT`
  - Upper bound for diagram enlargement in PDF output.
  - Lower it if diagrams are printing too large.
  - Raise it cautiously if diagrams look too small even when there is space.
- `PDF_PRINT_WIDE_DIAGRAM_ASPECT_RATIO`
  - Aspect-ratio threshold for promoting a diagram to landscape.
  - Lower it if wide diagrams are not rotating when they should.

- Raise it if too many diagrams are switching to landscape.
- `PDF_PRINT_WIDE_DIAGRAM_LANDSCAPE_GAIN`
  - Minimum width gain required before landscape is preferred.
  - Lower it if landscape should be chosen more aggressively.
  - Raise it if portrait should remain the default more often.
- `PDF_PRINT_PLANTUML_LANDSCAPE_ASPECT_RATIO`
  - Separate PlantUML-specific landscape trigger.
  - Useful because PlantUML often benefits from landscape earlier than Mermaid.
- `PDF_PRINT_HORIZONTAL_MARGIN_PX`
- `PDF_PRINT_VERTICAL_MARGIN_PX`
  - Effective printable margin budget used by the hidden print-layout solver.
  - Lower them to let diagrams use more page area.
  - Raise them if output feels cramped against the page edges.
- `PDF_PRINT_HEADING_TO_DIAGRAM_GAP_PX`
  - Vertical spacing reserved between a heading cluster and the diagram it governs.
  - Raise it if headed sections feel visually crowded.
  - Lower it if headed diagrams are spilling to the next page too eagerly.
- `PDF_PRINT_LAYOUT_SAFETY_PX`
  - Extra safety padding used to avoid borderline fits that clip in Chromium.
  - Lower it if diagrams are being pushed into spill mode too early.
  - Raise it if “just barely fits” cases are clipping or producing unstable output.

Heuristic guidance:

- If a diagram is being split across pages but you would accept smaller text, lower `MIN_PRINT_DIAGRAM_FONT_PT`.
- If a diagram is tiny but should clearly be landscape, first adjust `PDF_PRINT_WIDE_DIAGRAM_ASPECT_RATIO` or `PDF_PRINT_WIDE_DIAGRAM_LANDSCAPE_GAIN`.
- If headings are being orphaned above diagrams, leave the font knobs alone and look instead at the page-geometry and spacing constants.
- If many diagrams are near the right answer but consistently a little too big or too small, adjust margins before changing font caps.

Recommended workflow when tuning:

1. Change one constant at a time.
2. Regenerate the specific problematic PDF.
3. Recheck both the target document and one known-good document.
4. Prefer small adjustments; many of these settings interact.

The hidden PDF preflight logic consumes these constants through a Python-to-JS configuration handoff, so tuning values in `mdexplore.settings.json` is the intended way to adjust print behavior.

## Known TODOs

- Diagram interaction state restore is not yet reliable across document switches:
  - Mermaid zoom/pan state may not restore consistently when returning to a file.
  - PlantUML zoom/pan state may not restore consistently when returning to a file.
- This is tracked as an explicit TODO for future maintenance; core markdown rendering, search, highlighting, and PDF export remain functional.

## Multiple Views

- Click `Add View` to create another tab for the same currently previewed document.
- New tabs inherit the current view's top visible line/scroll position.
- By default, tab labels show the current top-most visible source line number for that tab.
- Tabs show a small left-side position bargraph indicating where that view sits within the document.
- Tabs use a fixed soft-pastel color sequence based on the order each view was opened.
- Tabs can be dragged to reorder without changing each tab's assigned color.
- Right-click a tab to assign a custom label (including spaces) up to 48 characters.
- If a longer label is entered, mdexplore truncates it to the first 48 characters.
- Entering a blank custom label restores the default dynamic line-number label for that tab.
- When a tab receives a custom label, mdexplore stores that tab's current scroll position and top visible source line as the tab's saved beginning.
- Right-click a custom-labeled tab to use `Return to beginning`, which jumps that tab back to the stored label-time location.
- Custom-labeled tabs also show a refresh icon beside the close button; clicking it resets that tab's saved beginning to the current scroll position/top line, the same way relabeling the tab does.
- Relabeling a custom-labeled tab with different text resets the saved beginning to the scroll position at the time of relabeling.
- When a new view is added and the palette wraps, mdexplore skips any color already used by open tabs.
- If you switch to another markdown file and later return in the same app run, that file's tabs restore with their prior order and selected tab.
- View-tab state also persists across app restarts in `.mdexplore-views.json` beside the document directory, keyed by markdown filename.
- For custom-labeled tabs, `.mdexplore-views.json` also persists the stored label-time beginning location used by `Return to beginning`.
- If custom labels make the tab strip too wide for the window, tab scrolling is enabled through the tab bar's built-in scroll buttons.
- Only documents with more than one saved view, or with a custom tab label, are written to `.mdexplore-views.json`; untouched single-view documents continue to use the default one-tab state.
- The tab strip is hidden when there is only one unlabeled default view.
- If only one view remains and it has a custom label, its tab stays visible so the custom label and `Return to beginning` action remain available.
- Closing that sole remaining custom-labeled tab clears the custom label and bookmark, then returns the document to the hidden default single-view state.
- Tabs are closeable with `X`; at least one tab is always kept open.
- Maximum views per document: `8`.
- Named-view left-gutter markers use the same restore path as selecting the corresponding tab, so marker navigation lands at the same saved location.

## Markdown Callouts

- mdexplore supports GitHub/Obsidian-style callouts written as blockquotes.
- Supported types: `NOTE`, `TIP`, `IMPORTANT`, `WARNING`, `CAUTION` (case-insensitive).
- `INFO` is accepted and styled as a `NOTE` callout.
- Custom titles are supported: `> [!WARNING] Custom title`.
- `+` / `-` markers are accepted in syntax, but callouts are rendered as non-collapsible boxes.

Example:

> [!NOTE]  
> This is a note callout with **markdown** content.

## Search and Match Highlighting

- Use `Search and highlight:` to match markdown files currently visible in the tree (root + expanded directories).
- While search is active, expand/collapse and directory/root scope changes automatically rerun the search against the newly visible set.
- Collapsed directories and all files beneath them are excluded from search and show no directory hit-count pill; collapsing a branch removes stale results immediately.
- File and recursive directory hit pills appear progressively as worker batches finish; symlink matches count beneath their visible alias path rather than beneath an unrelated resolved target path.
- Matching files are shown bold+italic in the tree.
- If filename terms match, filename text is rendered in yellow.
- The effective-root directory label is bold aqua when idle, and becomes dark orange with an appended hit-count pill when active search has hits under that scope.
- Press `Enter` in the search field to run search immediately (skip debounce).
- Clicking the `X` in the search field clears search text and removes match styling.
- Opening a matched file while search is active highlights matching text in yellow in the preview and scrolls to the first highlighted match.
- Preview scrollbar markers show where highlighted hits occur within the document; clicking a marker jumps to the nearest hit in that marker cluster.
- For `NEAR(...)` queries, preview highlighting is constrained to qualifying `NEAR` windows rather than every standalone occurrence of those terms elsewhere in the document.
- Non-quoted terms are case-insensitive.
- Double-quoted terms are case-insensitive and preserve spaces.
- Single-quoted terms are case-sensitive and preserve spaces.
- Only the quote character that opens a phrase closes it; the other quote character remains literal inside the phrase.
- Function-style operators accept both no-space and spaced forms before `( :`  
`NEAR(...)` / `NEAR (...)` , `OR(...)` / `OR (...)` , `AND(...)` / `AND (...)` ,  
and `NOT(...)` / `NOT (...)` .
- Legacy `CLOSE(...)` remains accepted for backward compatibility, but is normalized internally to canonical `NEAR(...)` .
- `AND(...)` , `OR(...)` , and `NEAR(...)` accept comma-delimited, space-delimited, or mixed argument lists.
- `AND(...)` and `OR(...)` are variadic and can take 2+ arguments.
- `NEAR(...)` requires 2+ terms and matches only when all terms appear within 50 words of each other in file content.
- `NEAR(...)` requires distinct qualifying occurrences for each listed term; one text start cannot satisfy two different required terms.
- For single-word `NEAR(...)` terms, proximity matching uses word boundaries, so `the` does not qualify via `The` in `They` .
- For `NEAR(...)` queries, the tree hit-count pill counts qualifying `NEAR` windows in the file, not the number of individual highlighted term spans.

Search examples:



```
joe
```

Matches `joe` , `JOE` , `JoE` , etc. (filename or content).

```
"Anne Smith"
```

Case-insensitive exact phrase match.

```
'Anne Smith'
```

Case-sensitive exact phrase match.

```
"Program Director's RAG"
```

Case-insensitive exact phrase match, with the apostrophe treated as literal text.

```
OR (Joe, "Fred")
```

Matches files containing either `Joe` (any case) or phrase `"Fred"` in any case.

```
NEAR(Fred "Anne Smith" Joe)
```

Matches only if all listed terms occur within 50 words of each other.

```
NEAR('The ', the)
```

Requires a distinct later/earlier standalone `the` ; the `The` in `They` does not qualify.

```
Joe "Anne Smith" NOT draft
```

Implicit `AND` : equivalent to `Joe AND "Anne Smith" AND NOT draft` .

```
AND(alpha beta gamma)
```

Requires all three terms to match (equivalent to `alpha AND beta AND gamma` ).

## PlantUML Local Configuration

The app renders PlantUML blocks locally with `plantuml.jar` .

- Default jar path: `/path/to/mdexplore/vendor/plantuml/plantuml.jar`
- Override jar path with `PLANTUML_JAR` :

```
PLANTUML_JAR=/path/to/plantuml.jar /path/to/mdexplore/mdexplore.sh
```

Behavior details:

- Markdown preview loads immediately with `PlantUML rendering...` placeholders.

- Each diagram is replaced automatically as soon as it completes.
- Diagram replacements are applied in-place so scroll position is preserved.
- PlantUML SVG data: URIs are BASE64-encoded through guarded helpers; if vendored `fastbase64` encode output is malformed, mdexplore automatically falls back to `pybase64` /stdlib encode so diagrams still render.
- Failed diagrams show `PlantUML render failed with error ...` plus detailed stderr context (including line number when PlantUML provides one).
- Diagram progress continues while you browse other files; returning shows completed progress.

## MathJax Local-First Configuration

Math rendering is local-first:

- mdexplore first tries a local `tex-svg.js` bundle (best rendering quality).
- If no local SVG bundle is found, it falls back to local `tex-mml-cthtml.js`.
- If no local bundle is found, it falls back to CDN.

Local lookup order:

1. `MDEXPLORE_MATHJAX_JS` (explicit file path)
2. `mathjax/es5/tex-svg.js` under the repo
3. `mathjax/tex-svg.js` under the repo
4. `assets/mathjax/es5/tex-svg.js` under the repo
5. `vendor/mathjax/es5/tex-svg.js` under the repo
6. System paths such as `/usr/share/javascript/mathjax/es5/tex-svg.js`
7. Local CHTML bundle paths ( `tex-mml-cthtml.js` ) as fallback

Example override:

```
MDEXPLORE_MATHJAX_JS=/absolute/path/to/tex-svg.js /path/to/mdexplore/mdexplore.sh
```

## Mermaid Local-First Configuration

Mermaid rendering is local-first:

- mdexplore first tries a local `mermaid.min.js` bundle.
- If no local bundle is found, it falls back to CDN.
- If `--mermaid-backend rust` is used, mdexplore uses `mmdr` instead of Mermaid JS for diagram SVG generation.

Local lookup order:

1. `MDEXPLORE_MERMAID_JS` (explicit file path)
2. `mermaid/mermaid.min.js` under the repo
3. `mermaid/dist/mermaid.min.js` under the repo
4. `assets/mermaid/mermaid.min.js` under the repo
5. `assets/mermaid/dist/mermaid.min.js` under the repo
6. `vendor/mermaid/mermaid.min.js` under the repo
7. `vendor/mermaid/dist/mermaid.min.js` under the repo
8. System paths such as `/usr/share/javascript/mermaid/mermaid.min.js`

Example override:

```
MDEXPLORE_MERMAID_JS=/absolute/path/to/mermaid.min.js /path/to/mdexplore/mdexplore.sh
```

Rust backend executable override:

```
MDEXPLORE_MERMAID_RS_BIN=/absolute/path/to/mmdr /path/to/mdexplore/mdexplore.sh --mermaid-backend rust
```

## Markdown Engine and BASE64 Performance Knobs

- Markdown engine selection:
  - `MDEXPLORE_MARKDOWN_ENGINE=cmark` (default): uses `cmarkgfm`.
  - `MDEXPLORE_MARKDOWN_ENGINE=markdown-it`: forces `markdown-it-py`.
  - If `cmarkgfm` is unavailable at runtime, mdexplore automatically falls back to `markdown-it-py`.

Example:

```
MDEXPLORE_MARKDOWN_ENGINE=markdown-it /path/to/mdexplore/mdexplore.sh
```

- BASE64 image worker threads (preview data-URI materialization + copy-time prefetch):
  - `MDEXPLORE_BASE64_IMAGE_THREADS=<N>`
  - Default is CPU-based (`max(2, min(24, cpu_count * 2))`).

Example:

```
MDEXPLORE_BASE64_IMAGE_THREADS=24 /path/to/mdexplore/mdexplore.sh
```

- BASE64 codec path:
  - Decode path prefers `pybase64`, with adaptive routing to vendored `vendor/fastbase64` for payload-size sweet spots.
  - Encode path may use vendored `fastbase64`, but output is validated (length/padding/alphabet). Invalid vendor output is discarded and mdexplore falls back to `pybase64/stdlib` automatically.
  - Adaptive benchmark state persists in `fastbase64-adaptive.json`.
  - Optional vendor-vs-fallback benchmark telemetry is written to `fastbase64-benchmark.log` when the debug dual-path switch is enabled in `mdexplore_app/fast_base64.py`.
  - If accelerated backends are unavailable, helpers fall back to Python stdlib `base64`.

## Debug Logging

- `mdexplore.py --debug` enables verbose debug logging to `mdexplore.log` in the project directory.
- `mdexplore.sh` exposes the same behavior through the `DEBUG_MODE` variable near the top of the script. Set `DEBUG_MODE=true` to pass `--debug` into the app.
- When debug logging is disabled, `mdexplore.log` is not written.
- The application debug log trims itself to the most recent `10,000` lines.
- The non-interactive launcher log remains separate at `~/.cache/mdexplore/launcher.log` and trims to `1,000` lines.

## Render and Rules Maps

- See [DEVELOPERS-AGENTS.md](#) for the consolidated deep maps:
  - formal behavior/rule hierarchy,
  - render/caching and PDF/export control flow,
  - decision tables and invariants,
  - agent-facing maintenance workflow.
- See [UML.md](#) for subsystem-level architecture, class relationships, and activity diagrams.

## Project Structure

```
mdexplore.py          # Qt application, file tree, renderer integration
mdexplore.sh          # launcher (venv create/install/run)
setup-mdexplore.sh    # full bootstrap script (venv/assets/mmdr)
mdexplore.desktop.sample # sample desktop launcher entry for user customization
mdexplore_app/        # extracted support modules for search, tree, tabs, icons, PDF,
runtime, template/JS assets, and workers
tests/                # headless/unit regressions for preview, search, template assets, and
tab layout
assets/js/            # externalized preview/PDF JavaScript templates loaded into a startup
registry at startup
assets/js/pdf/        # PDF export/preflight/restore JavaScript templates
assets/js/preview/    # preview search/highlight/selection/context-menu JavaScript
templates
assets/templates/    # externalized HTML preview document templates loaded into a startup
registry
assets/templates/preview/ # preview HTML shell/template assets
assets/ui/            # local UI icons/font assets used by tree, tabs, copy controls, and
search pills
mdexplor-icon.png     # primary app icon asset kept at repo root for desktop launchers
requirements.txt       # Python runtime dependencies
README.md             # project docs
DEVELOPERS-AGENTS.md  # developer and coding-agent maintenance guide
UML.md               # subsystem-level architecture/class/activity diagrams
LICENSE               # MIT license
```

## Development

Basic local checks:

```
bash -n mdexplore.sh
python3 -m py_compile mdexplore.py
```

## Troubleshooting

If you see `ModuleNotFoundError: No module named 'PySide6.QtWebEngineWidgets'`:

- Re-run `mdexplore.sh`; if runtime import verification is triggered and fails, it auto-reinstalls dependencies.

- If verification was previously stamped but the venv later drifted, force a recheck by removing:
  - `/path/to/mdexplore/.venv/.runtime-import-check.sha256`
 then run:
  - `/path/to/mdexplore/mdexplore.sh`
- If it still fails, run:
  - `rm -rf /path/to/mdexplore/.venv`
  - `/path/to/mdexplore/mdexplore.sh`

If `Edit` does nothing:

- Ensure MarkText is installed.
- Verify `/usr/bin/marktext` exists and is executable.

If the dock/menu still shows an old app icon:

- Ensure your launcher contains:
  - `Icon=/absolute/path/to/mdexplor-icon.png`
  - `StartupWMClass=mdexplore`
- Refresh desktop entries:
  - `update-desktop-database ~/.local/share/applications`
- Remove old pinned `mdexplore` favorites from the dock and pin it again.
- Log out/in if the shell still shows the previous cached icon.

If running the launcher appears to do nothing:

- Watch terminal output; the launcher now prints setup/launch status.
- First dependency install can take time because Qt packages are large.
- Ensure you are in a GUI session ( `DISPLAY` or `WAYLAND_DISPLAY` must be set).
- For desktop/dock launches ( `Terminal=false` ), check launcher log:
  - `~/.cache/mdexplore/launcher.log`
  - log is auto-trimmed to the most recent 1000 lines

## Security Notes

- Markdown HTML is enabled ( `html=True` ), so preview untrusted Markdown with care.
- Mermaid uses local-first loading with CDN fallback.
- MathJax uses local-first loading with CDN fallback.

## Contributing

1. Keep changes focused and small.
2. Run syntax checks before submitting.
3. Update docs ( `README.md` , `DEVELOPERS-AGENTS.md` , and `UML.md` when relevant) when behavior changes.

## License

This project is licensed under the MIT License. See `LICENSE` .